



AI That Learns, Thinks, and Acts: The Next Frontier of Generative AI

Richard Shan , Consultant

This article dives into the recent transformation of artificial intelligence (AI) systems from passive content generators into autonomous decision makers and executors through the advancements of self-improving retrieval-augmented generation, generative AI-native agents, and large action models, which highlight their potential to automate workflows and enhance real-world applicability.

Generative artificial intelligence (GenAI) has rapidly evolved over the past few years,¹ fundamentally transforming industries ranging from software development to scientific research.² The introduction of large language models (LLMs) such as OpenAI's GPT-4o, Google's Gemini, and Meta's Llama has propelled AI into mainstream applications,³ enabling humanlike text generation, image synthesis, and multimodal reasoning.⁴ However, the next frontier of AI development goes far beyond content

generation,⁵ toward systems that can self-improve; act autonomously; and execute complex, real-world tasks.⁶

In 2025, three key advancements are set to redefine the landscape of GenAI, as illustrated in [Figure 1](#):

1. *Self-improving retrieval-augmented generation (RAG)*: AI systems that adapt and refine their knowledge in real time by integrating user feedback and continuous learning pipelines.
2. *GenAI-native agents*: Autonomous AI agents capable of multistep reasoning, real-world execution, and dynamic problem solving without human intervention.

3. *Large action models (LAMs)*: The shift from AI generating text to AI executing actions in digital and physical environments, bridging the gap between decision making and real-world impact.

These advancements represent a fundamental shift from passive AI assistants to proactive, agentic systems capable of learning from their environments, making autonomous decisions, and executing tasks with minimal human oversight. The implications of these developments extend across scientific research, healthcare, finance, and robotics, potentially unlocking new economic opportunities while also posing novel ethical and security challenges.

This article explores the technological foundations, emerging applications, and societal implications of self-improving RAG, GenAI-native agents, and LAMs in 2025. It also examines the potential impacts, challenges, and risks associated with these advancements and outlines key actions required to ensure responsible AI development. As AI continues to evolve, understanding the trajectory of GenAI in 2025 will be critical for researchers, industry leaders, and policy makers alike.

THE RISE OF SELF-IMPROVING RAG

The evolution of RAG

RAG has become an essential enhancement for LLMs,⁷ allowing them to retrieve relevant knowledge from external sources rather than relying solely on their pretrained parameters.⁸ Unlike traditional GenAI models, which are limited by the static nature of their training data, RAG

improves the accuracy and contextual awareness of AI-generated responses by dynamically integrating updated information. This approach has been particularly valuable in applications requiring domain-specific expertise, such as legal research, scientific discovery, and financial analysis.

However, early implementations of RAG were largely static, requiring manual intervention to update retrieval strategies or integrate new knowledge. While these systems significantly improved factual accuracy compared with standard GenAI models, they still lacked the ability to autonomously refine their own retrieval processes over time. As AI continues to advance in 2025, RAG is transitioning into a self-improving paradigm, where models autonomously adapt their retrieval strategies, optimize response generation, and enhance their knowledge in real time.

Self-improving RAG: A new frontier

The next generation of RAG systems is moving beyond simple retrieval to incorporate continuous learning and adaptive feedback loops. These self-improving models are designed to analyze user interactions, detect inaccuracies,

and refine their responses dynamically, reducing reliance on manual updates. By leveraging automated feedback mechanisms, AI systems can identify incorrect or outdated information and adjust their retrieval strategies accordingly. This self-correction process enhances factual accuracy while minimizing the need for human intervention.

Another key advancement in self-improving RAG is the integration of real-time knowledge ingestion pipelines. Rather than depending on static vector databases, modern RAG architectures can continuously pull data from trusted sources, including scientific repositories, enterprise knowledge bases, and regulatory updates. This ensures that AI-generated content remains up to date, even in rapidly evolving domains. Additionally, reinforcement learning techniques allow these models to refine their retrieval methods based on user feedback, enabling continuous optimization of search relevance and ranking mechanisms.

The development of fine-tuned vector search and adaptive indexing further enhances the efficiency of self-improving RAG. Advances in semantic search and hybrid retrieval techniques enable AI models to retrieve and rank documents with greater




FIGURE 1. GenAI prediction. RAG: retrieval-augmented generation.

precision. This results in more contextually relevant responses, reducing the risk of AI-generated misinformation. Together, these innovations mark a significant shift toward AI systems that do not merely generate responses but actively refine and evolve their understanding of information over time.

Key applications and industry impact

The introduction of self-improving RAG will have profound implications across multiple industries, particularly in sectors that rely on dynamic knowledge integration. In healthcare and scientific research, AI-powered medical assistants will leverage continuously updated medical literature to enhance disease diagnosis, treatment recommendations, and clinical decision support. By integrating real-time patient feedback and emerging medical studies, these AI systems will improve diagnostic accuracy and optimize personalized treatment plans.

In financial services, self-improving AI models will revolutionize risk assessment and fraud detection by dynamically processing market trends, regulatory changes, and financial reports. The ability to continuously ingest and analyze financial data will allow AI-driven systems to provide more accurate investment recommendations, predict economic fluctuations, and identify fraudulent transactions with greater precision. Similarly, corporate knowledge management will be transformed by AI-powered systems that automatically curate, summarize, and update enterprise-wide information. These intelligent knowledge bases will reduce the burden of manual documentation, ensuring that employees have access to the most relevant and up-to-date information at all times.

In the legal and regulatory domains, self-improving RAG will empower AI-driven legal assistants to stay updated with constantly evolving case law and regulations. Legal professionals will be able to access the most relevant precedents and statutory changes without manually searching through vast repositories of legal documents. Moreover, AI-powered research agents will redefine the way scientific literature is reviewed and synthesized. These systems will autonomously scan academic papers, generate hypotheses, and suggest novel research directions, accelerating innovation in various fields.

Challenges and limitations

Despite its transformative potential, self-improving RAG presents several technological and ethical challenges. One of the primary concerns is the risk of AI-generated misinformation and hallucination. While retrieval mechanisms enhance factual grounding, the integration of unreliable sources or improper ranking of retrieved documents could introduce incorrect information into AI-generated responses. Addressing this issue requires the development of robust fact-verification pipelines and automated credibility assessments to ensure that the retrieved knowledge is both accurate and reliable.

Another significant challenge lies in the scalability of real-time knowledge updates. Maintaining high-performance indexing, retrieval latency, and efficient knowledge ingestion at scale requires significant computational resources. Enterprises deploying self-improving RAG models must invest in infrastructure optimization to balance efficiency with accuracy. Additionally, the security of these AI systems remains a concern as adversarial attacks and data poisoning strategies

could manipulate retrieval mechanisms, introducing bias or misinformation into AI-generated outputs.

Privacy and data governance also present critical challenges, particularly in sectors such as healthcare and finance, where sensitive information must be protected. The ability of RAG systems to pull real-time data from multiple sources raises concerns about compliance with data protection regulations. Implementing secure retrieval mechanisms and privacy-preserving techniques will be essential in ensuring that self-improving AI adheres to ethical and legal standards.

The road ahead for self-improving RAG

As self-improving RAG continues to evolve, several research directions will be critical in shaping its development. Enhancing AI explainability will be a key priority, ensuring that retrieval processes remain transparent and interpretable. Fact-verification mechanisms must be integrated to minimize misinformation risks, while optimization efforts should focus on improving the scalability and efficiency of real-time knowledge pipelines.

Equally important is the establishment of ethical and regulatory safeguards to govern AI-driven knowledge retrieval. Defining best practices for AI-generated content, data security, and bias mitigation will be essential in ensuring responsible deployment. Collaborative efforts among AI researchers, policy makers, and industry leaders will play a crucial role in setting guidelines for the ethical use of self-improving AI.

By addressing these challenges, self-improving RAG will become a foundational component of next-generation AI systems. These advancements will enable AI models to evolve dynamically,

enhancing their ability to provide accurate, contextually relevant, and continuously updated knowledge across a wide range of applications. As AI continues to move beyond static knowledge retrieval, self-improving RAG represents a significant step toward truly adaptive and intelligent AI systems capable of learning and refining their understanding in real time.

GenAI-NATIVE AGENTS: THE NEXT EVOLUTION OF AI AUTONOMY

The shift from static AI to autonomous agents

The progression of AI from static language models to autonomous, decision-making systems represents one of the most significant technological leaps in 2025. GenAI-native agents are AI systems that move beyond simple text-based interactions, becoming proactive problem solvers capable of planning, reasoning, and executing complex, real-world tasks. Unlike conventional LLMs, which rely on pretrained knowledge and user prompts, GenAI-native agents are self-directed, continuously learning, and capable of multistep reasoning across diverse domains.

This transition is fueled by advancements in long-term memory architectures, multiagent collaboration, and real-world execution frameworks. AI agents are now able to autonomously retrieve knowledge,⁹ adapt to new environments, and make real-time decisions without direct human intervention.¹⁰ In 2025, these agents are expected to transform industries such as software development, research, cybersecurity, and automation, enabling AI to function as an independent digital workforce that can assist professionals across disciplines.

Core capabilities of GenAI-native agents

Unlike traditional LLM-based chatbots, which generate responses in a reactive manner, GenAI-native agents exhibit a set of advanced cognitive and execution capabilities that make them more autonomous and efficient.

with application programming interfaces (APIs), databases, robotic systems, and digital platforms. They can autonomously write and execute code, interact with software tools, manage cloud-based systems, and even control Internet of Things devices. As AI agents become embedded within enterprise software

DEFINING BEST PRACTICES FOR AI-GENERATED CONTENT, DATA SECURITY, AND BIAS MITIGATION WILL BE ESSENTIAL IN ENSURING RESPONSIBLE DEPLOYMENT.

One of the defining features of these agents is their self-directed problem-solving ability. Rather than responding to individual prompts, GenAI-native agents break down complex problems into actionable steps, plan solutions, and autonomously execute tasks across multiple systems. This capability enables AI-powered project managers, autonomous research assistants, and AI-driven software engineers to operate with minimal human intervention.

Another key aspect is multistep reasoning and planning. AI agents can contextually track long-term objectives, dynamically adjusting their strategies based on evolving information. Unlike traditional models that operate in single-turn conversations, GenAI-native agents maintain persistent memory, allowing them to refine their approach over time. This makes them particularly useful in applications such as automated scientific research, complex business workflows, and AI-driven diagnostics in healthcare.

Additionally, these agents are capable of real-world execution by interfacing

ecosystems, they will fundamentally redefine automation by reducing the need for manual intervention in data analysis, cybersecurity threat response, and business decision making.

Key applications and industry impact

The rise of GenAI-native agents is expected to reshape multiple industries, offering new ways to automate complex workflows, optimize research, and improve decision making. One of the most immediate impacts will be seen in software development, where AI-powered coding agents can autonomously write, debug, and optimize code. Tools like OpenAI's Code Interpreter, Meta's LLM-driven development assistants, and experimental AI software engineers are already demonstrating the potential for AI to function as an autonomous codeveloper, accelerating software production cycles.

In business process automation, AI-driven workflow agents will revolutionize customer service, financial forecasting, and enterprise management

by autonomously processing documents, managing schedules, and analyzing large datasets for actionable insights. These AI-driven automation systems will be capable of handling end-to-end processes, reducing the need for manual intervention in routine administrative tasks.

Scientific research will also benefit from AI-native agents acting as autonomous research assistants. These systems will be capable of conducting literature reviews, formulating hypotheses, and even running simulations to test scientific theories. For example, AI-powered chemistry research agents will be able to propose novel compounds for drug discovery, while AI-driven physics models can assist in solving complex equations related to quantum mechanics and material science.

In the cybersecurity domain, AI-native agents will function as autonomous threat response systems, continuously scanning for vulnerabilities, detecting anomalies, and executing preemptive countermeasures against cyberthreats. These AI-driven security models will enhance adaptive defense strategies for enterprises, government agencies, and financial institutions.

Furthermore, the integration of GenAI-native agents into autonomous vehicles, robotics, and smart infrastructure will pave the way for AI-driven logistics, warehouse management, and transportation networks. AI-native agents will manage autonomous supply chain operations, predict equipment failures in industrial settings, and coordinate multiagent robotic systems in warehouses.

Challenges and limitations

Despite their transformative potential, GenAI-native agents face several technological, ethical, and operational

challenges that must be addressed before widespread adoption. One of the most significant hurdles is the challenge of long-term memory and contextual awareness. While current AI models exhibit impressive short-term reasoning capabilities, maintaining consistent memory across extended interactions remains difficult. Researchers are exploring novel architectures, such as hierarchical memory models and reinforcement learning-enhanced planning, to improve AI's ability to track objectives over time.

Another critical concern is the reliability and robustness of AI decision making. Autonomous agents must operate in dynamic, unpredictable environments where unexpected inputs or errors could lead to flawed decision making. Ensuring robust AI reasoning, safety checks, and human-AI oversight mechanisms is crucial to preventing unintended actions in high-risk scenarios, such as healthcare diagnostics, financial transactions, and autonomous robotics.

Security and adversarial manipulation are also pressing concerns. As AI-native agents become more integrated into enterprise environments, they will become attractive targets for cyberattacks and adversarial exploits. Ensuring that AI agents can defend against prompt injection attacks, adversarial perturbations, and data manipulation tactics will require robust security architectures and continuous monitoring.

Finally, ethical considerations must be addressed, particularly regarding accountability and transparency. As AI agents gain autonomy in executing decisions, determining responsibility for AI-driven actions becomes more complex. Policy makers and researchers must work together to establish regulatory frameworks that ensure responsible AI deployment while maintaining human oversight.

The road ahead for GenAI-native agents

The development of trustworthy, scalable, and efficient GenAI-native agents will require continued advancements in long-term memory architectures, interpretability, and safety protocols. AI developers must focus on enhancing AI's ability to explain its reasoning processes, ensuring that decision making remains transparent and auditable. Additionally, incorporating human-in-the-loop oversight will be essential in mission-critical applications where AI decisions have significant consequences.

As AI-native agents continue to evolve, collaboration among AI researchers, policy makers, and industry stakeholders will be critical in defining ethical guidelines and best practices for AI autonomy. The successful deployment of these agents will mark a significant shift toward AI-powered automation across industries, where AI is not just a tool for assisting humans but an active participant in decision making, problem solving, and real-world execution.

By addressing the technological and ethical challenges associated with AI-native agents, researchers can unlock a future where AI systems act as intelligent collaborators, augmenting human capabilities and driving innovation across multiple domains. In 2025, the emergence of these AI-driven entities will not only accelerate digital transformation but also redefine the role of artificial intelligence as a truly autonomous force in shaping the future of technology.

LAMs: AI THAT GOES BEYOND WORDS

From text generation to action execution

GenAI has predominantly been associated with content creation, such as

text, images, audio, and code. However, the next frontier in AI advancement is action-based AI, where models can not only generate content but also execute complex tasks autonomously. LAMs represent this paradigm shift, transitioning AI from passive knowledge generators to decision-making and action-performing agents in digital and physical environments.

Unlike traditional LLMs that provide recommendations or explanations, LAMs have the ability to interact with external systems, trigger workflows, and perform sequential, goal-oriented actions. These models go beyond standard reinforcement learning frameworks, integrating causal reasoning, environmental interaction, and real-world task execution. By bridging the gap between generative intelligence and operational automation, LAMs will play a foundational role in industries ranging from robotics to software automation and enterprise-level decision-making systems.

What makes LAMs different?

The defining characteristic of LAMs is their ability to execute tasks autonomously across multiple domains, requiring them to process information, plan actions, and engage in real-world decision making. Unlike standard AI models, which primarily function in textual and virtual spaces, LAMs integrate three fundamental capabilities:

- › **Task-oriented orchestration:** LAMs are designed to complete pre-defined goals, breaking down complex objectives into multi-step sequences. These AI models can interface with APIs, robotic hardware, databases, and workflow automation tools to carry out real-world operations.

- › **Environmental adaptability:** Unlike conventional models, which rely on static inputs, LAMs can adapt to new information dynamically, refining their execution strategies based on real-time feedback. This adaptability allows them to operate in nondeterministic environments, where uncertainties must be accounted for.

- › **Causal and strategic planning:** LAMs incorporate long-term planning, causal inference, and counterfactual reasoning to anticipate the consequences of their actions. They can execute complex strategies, making them invaluable in fields like supply chain logistics, financial forecasting, and medical treatment planning.

These AI models leverage a combination of reinforcement learning, large-scale retrieval mechanisms, and real-time decision trees to assess their environment, plan actions, and execute tasks autonomously. The result is a new breed of AI that not only understands problems but also actively works toward solutions.

Key innovations driving LAMs

The rapid advancements in LLM-based automation, AI-enhanced robotics, and multiagent systems have paved the way for the development of LAMs. Several critical innovations are enabling this shift.

Multimodal action execution. LAMs are built to process and execute tasks across multiple modalities, including text, vision, speech, and motor control. By integrating LLMs with vision models and reinforcement learning systems, AI can interact with digital interfaces, physical hardware, and autonomous agents simultaneously.

This capability is crucial for applications in robotics, logistics, and autonomous vehicle navigation.

Fine-tuned causal inference. One of the most challenging aspects of AI-driven automation is the ability to predict outcomes based on causal relationships. Unlike conventional AI models that react to inputs in a sequential manner, LAMs utilize causal learning techniques to infer how actions impact future outcomes, allowing them to make strategic, long-term decisions. This is particularly valuable in financial market analysis, urban infrastructure management, and complex engineering systems.

AI-driven workflow automation. With the integration of API-based execution capabilities, LAMs can automate enterprise workflows, manage IT operations, and execute software deployment pipelines. This transforms AI from a recommendation engine into a fully operational automation layer, optimizing business processes with minimal human oversight. These models will enable self-managing cloud platforms, AI-driven cybersecurity enforcement, and AI-operated digital governance systems.

Impact on industries and real-world applications

LAMs will revolutionize multiple industries by transforming AI into an action-oriented decision-making framework. Several key areas stand to benefit from this evolution.

Manufacturing and robotics. In industrial automation, smart manufacturing, and robotics, LAMs will act as intelligent controllers for robotic arms, supply chain management, and warehouse logistics. AI-driven predictive maintenance systems will autonomously

identify mechanical failures before they occur, preventing costly downtime.¹¹

Healthcare and AI-driven medical assistance. In medical diagnostics and treatment planning, LAMs will assist doctors by autonomously analyzing patient history, running diagnostic tests, and optimizing treatment protocols.

Autonomous vehicles and smart infrastructure. LAMs will play a pivotal role in autonomous vehicle navigation, drone operations, and smart city infrastructure. AI-driven transportation systems will manage real-time traffic routing, predict congestion patterns, and optimize urban planning decisions using autonomous AI coordination.¹⁵

becomes a pressing issue. Developing regulatory frameworks that establish clear guidelines on liability and oversight will be crucial for the safe deployment of LAMs.

The road ahead for LAMs

To fully realize the potential of LAMs, further research is required in causal inference, real-time adaptability, and safety validation. AI developers must integrate explainable decision-making protocols, ensuring that AI-driven actions remain transparent and auditable.

Collaboration among AI researchers, industry leaders, and policy makers will be necessary to establish best practices, regulatory safeguards, and ethical deployment standards for LAM-driven automation. Ensuring that LAMs operate within ethical, legal, and operational constraints will be key to their successful integration into critical industries.

As AI progresses from generating knowledge to executing actions, the emergence of LAMs will redefine how machines interact with the world. These AI-driven systems will not only automate workflows and execute decisions but also adapt, learn, and improve in real time, marking a significant step toward fully autonomous AI-powered ecosystems.

RISKS TO PREDICTION

Unforeseen breakthroughs or limitations

While the current trajectory of self-improving RAG, GenAI-native agents, and LAMs appears promising, the precise evolution of AI remains inherently unpredictable. Unexpected breakthroughs in alternative AI architectures, such as neuromorphic computing, quantum AI, or hybrid symbolic-reasoning systems, could significantly alter the

ENSURING THAT LAMs OPERATE WITHIN
ETHICAL, LEGAL, AND OPERATIONAL
CONSTRAINTS WILL BE KEY TO THEIR
SUCCESSFUL INTEGRATION INTO
CRITICAL INDUSTRIES.

In robotic-assisted surgery, AI will enhance precision and adaptability, allowing real-time adjustments based on patient-specific factors.¹²

AI-powered finance and algorithmic trading. Financial institutions will deploy LAMs to execute high-frequency trading strategies, dynamic risk assessments, and portfolio optimizations. Unlike traditional AI models that provide market insights, LAMs will autonomously adjust trading parameters, respond to market shifts, and execute real-time trades, significantly enhancing financial decision making.¹³

Cybersecurity and autonomous threat detection. In cybersecurity, LAMs will actively defend networks by detecting, analyzing, and neutralizing cyberthreats before human intervention is required. These AI systems will autonomously apply security patches, mitigate attack vectors, and conduct real-time forensic analysis on potential breaches.¹⁴

Challenges and limitations

Despite their transformative potential, LAMs present significant technical and ethical challenges. One of the most critical limitations is ensuring the reliability of AI-driven actions in high-stakes environments. Errors in execution could lead to catastrophic failures in healthcare, finance, and autonomous systems, making rigorous validation, testing, and safety protocols essential.

Another major challenge is handling uncertainty in dynamic environments. LAMs operating in nondeterministic settings must be capable of adapting their strategies when unexpected variables arise. This requires advances in adaptive reinforcement learning, human-AI collaboration, and real-time simulation environments.

Ethical concerns also emerge regarding AI accountability and decision-making authority. As LAMs gain autonomy in executing real-world actions, determining responsibility for AI-driven errors or unintended consequences

developmental path. However, rather than rendering existing approaches obsolete, these innovations could integrate with and enhance current models, providing complementary strengths and creating new opportunities for innovation. Proactively exploring such synergies will help ensure that ongoing advancements in AI remain robust, adaptable, and aligned with future technological possibilities.

Conversely, technical limitations may hinder progress. The increasing size and complexity of AI models demand exponential computing power, and without significant advances in energy-efficient AI hardware or distributed learning techniques, continued growth may become economically or environmentally unsustainable. If research into scalable and resource-efficient AI architectures does not keep pace, the widespread adoption of highly autonomous AI systems could face major roadblocks.

Regulatory and ethical roadblocks

The regulatory landscape for AI governance is evolving rapidly, and its trajectory could significantly impact how AI autonomy, decision making, and automation are deployed in the coming years. As AI systems become more powerful and capable of independent actions, governments and policy makers may enforce stricter regulations to limit the risks associated with misuse, safety failures, and ethical concerns.

For instance, global AI governance frameworks, such as the European Union AI Act and proposed U.S. regulations, are already outlining compliance requirements for high-risk AI applications. If regulatory scrutiny increases beyond anticipated levels, AI developers and organizations may be forced to slow down deployment

or implement additional safety layers that could hinder efficiency. Alternatively, fragmented AI regulations across different countries may create barriers to global AI standardization, leading to regional disparities in AI adoption and accessibility.

Another major concern is ethical resistance to AI decision making in high-stakes fields. While AI-native agents and LAMs could significantly enhance automation in healthcare, law enforcement, and governance, societal hesitation around machine-driven judgments and algorithmic accountability may delay widespread adoption. Ensuring explainability, fairness, and public trust in AI autonomy will be critical in overcoming these regulatory and ethical hurdles.

Public and enterprise adoption uncertainty

Even if technological advancements and regulatory frameworks align, the success of AI-native autonomy ultimately depends on enterprise and public adoption. Despite AI's rapid evolution, human skepticism, resistance to automation, and concerns about workforce displacement could slow adoption across industries.

Many organizations, particularly in traditional sectors such as law, finance, and government, may be reluctant to replace human decision makers with AI-native agents, fearing liability issues, ethical concerns, or reputational damage. Additionally, while AI automation can enhance efficiency, companies may struggle with integrating AI-driven decision making into existing workflows, particularly if employees resist AI-assisted job augmentation.

From a public perspective, there is growing concern about AI's influence on job security, misinformation, and decision transparency. While

AI-native systems may be more efficient and objective, they may also be perceived as lacking human empathy, creativity, and contextual understanding, factors that are essential in fields such as education, mental healthcare, and creative industries. Public distrust in AI's ability to make nuanced, socially responsible decisions could lead to pushback against widespread AI deployment, forcing organizations to implement human-AI hybrid solutions rather than fully autonomous systems.

Cybersecurity and AI vulnerabilities

As AI systems gain more autonomy and decision-making capabilities, the risk of AI-powered cyberthreats increases significantly. While AI is already being used in cybersecurity defense, it is also being weaponized in automated hacking, misinformation campaigns, and adversarial AI attacks.

One of the greatest risks is that self-improving RAG models could be manipulated through data poisoning attacks, where malicious actors inject biased, misleading, or harmful information into AI retrieval mechanisms. Since self-improving AI systems autonomously refine their knowledge, a compromised RAG pipeline could result in widespread AI misinformation, manipulated decision making, or algorithmic exploitation.

Similarly, AI-native agents and LAMs present new attack surfaces as these systems directly interact with external applications, APIs, and digital environments. If adversaries gain access to AI-powered automation pipelines, they could alter AI-driven financial transactions, disrupt enterprise automation, or manipulate AI-assisted governance systems. The ability for

bad actors to influence AI's operational decision making could pose severe risks to national security, economic stability, and digital trust.

To mitigate these risks, organizations must implement secure AI architectures, real-time adversarial defense mechanisms, and continuous monitoring of AI-driven decision making. Cybersecurity must evolve alongside AI, ensuring that AI-native automation remains resistant to manipulation, tampering, or unauthorized control.

The limits of AI reasoning and adaptability

While self-improving RAG, GenAI-native agents, and LAMs aim to enhance AI autonomy, fundamental questions remain about whether AI can truly replicate humanlike reasoning, intuition, and adaptability. Despite advances in causal learning, multistep reasoning, and action-based AI, current AI models still struggle with abstract thinking, common sense reasoning, and context-based decision making in ambiguous or novel situations.

For example, AI systems trained on historical data struggle with out-of-distribution learning, meaning they may fail to generalize effectively when encountering entirely new scenarios. While reinforcement learning can improve adaptability, AI still lacks the true cognitive flexibility of human intelligence, particularly in areas requiring ethical judgment, emotional intelligence, and strategic foresight.

Additionally, multiagent AI collaboration introduces further unpredictability. When multiple AI agents coordinate in real-time environments, they may develop unintended emergent behaviors, potentially leading to unexpected or chaotic system outputs. Ensuring alignment among AI goals,

constraints, and human values will require advanced AI alignment research and strict operational oversight.

The road ahead for AI predictions

Despite these risks, the development of self-improving RAG, GenAI-native agents, and LAMs is likely to continue advancing rapidly, driven by enterprise demand, technological breakthroughs, and ongoing research into AI reasoning and autonomy. However, the success of these AI-driven innovations will depend on how well researchers, organizations, and policy makers address the challenges of AI security, governance, and societal integration.

By combining AI advancements with strong regulatory frameworks, ethical oversight, and robust cybersecurity measures, the AI industry can ensure that autonomous AI remains a force for innovation rather than disruption. However, failure to properly secure, regulate, and monitor AI autonomy could result in unintended consequences that undermine trust, stability, and ethical deployment.

The future of AI is not just a question of technological capability; it is a question of how responsibly and strategically AI autonomy is deployed. As AI systems become increasingly embedded in real-world automation, governance, and decision making, ensuring that these technologies align with human values, safety standards, and long-term ethical considerations will be the defining challenge of AI's next evolution.

Call to action

For AI researchers and developers: Advancing safe and responsible AI.

The continued development of self-improving RAG, GenAI-native agents, and LAMs depends on researchers and

developers prioritizing transparency, safety, and scalability. As AI systems evolve from passive tools to autonomous agents capable of self-improvement and real-world execution, it is critical to implement robust evaluation mechanisms, interpretability frameworks, and ethical safeguards.

To advance AI safety, researchers must focus on enhancing long-term memory architectures, minimizing bias, and ensuring AI-generated knowledge remains accurate and verifiable. The implementation of hybrid AI architectures, combining symbolic reasoning with deep learning approaches, could provide more reliable and interpretable AI systems. Furthermore, fact-verification pipelines and adversarial robustness training should be integrated into self-improving RAG systems to prevent misinformation, biased retrievals, or manipulation of AI-driven insights.

In the case of GenAI-native agents and LAMs, AI developers must implement safety constraints, action validation checks, and human-in-the-loop oversight models to prevent unintended or dangerous behaviors in real-world execution. Additionally, advancements in explainable AI are essential to increase transparency in AI decision making, ensuring users can audit, challenge, and modify AI-generated outputs when necessary.

For enterprises and policy makers: Ensuring AI governance and ethical deployment.

As AI systems become increasingly embedded in business operations, decision-making processes, and government frameworks, enterprises and policy makers must take proactive steps to ensure AI deployment aligns with legal, ethical, and security standards. This includes developing clear regulatory

frameworks that balance innovation with risk mitigation, ensuring that AI-driven decisions remain accountable and transparent.

Governments should work closely with AI researchers and industry leaders to establish standardized compliance guidelines for AI deployment, particularly in high-risk domains such as healthcare, finance, and autonomous infrastructure management. The introduction of AI auditing mechanisms, real-time monitoring frameworks, and bias detection algorithms will be critical in minimizing unintended consequences and ensuring fairness in AI-driven automation.

From an enterprise perspective, companies adopting AI-native automation must prioritize AI ethics, security, and compliance. Developing internal AI governance policies, appointing AI ethics committees, and investing in AI risk assessment tools will be essential to ensuring responsible AI integration. Organizations must also educate employees on AI capabilities, limitations, and best practices, fostering a culture of human-AI collaboration rather than fear of displacement.

For society: Preparing for an AI-augmented future. As AI continues to reshape industries, economies, and daily life, society must adapt to the realities of AI-driven transformation. One of the most immediate challenges is workforce displacement as AI-native automation is expected to eliminate certain job roles while creating new opportunities in AI governance, system oversight, and AI-human collaboration. Governments, educational institutions, and industry leaders must prioritize large-scale reskilling initiatives, ensuring that workers can transition into AI-driven job markets.

Public awareness campaigns should focus on improving AI literacy, helping individuals understand AI's capabilities, risks, and ethical considerations. Addressing public concerns about AI decision making, bias, and accountability will be crucial in ensuring widespread acceptance and responsible adoption of AI technologies.

At a societal level, it is important to establish ethical norms around AI autonomy, ensuring that AI-driven decisions remain aligned with human values and democratic principles. Civil organizations and research institutions should advocate for AI fairness, inclusivity, and accessibility, ensuring that AI benefits all segments of society rather than exacerbating existing inequalities.

The vision for 2030: The next phase of AI evolution. While 2025 marks a turning point in AI autonomy, the long-term vision extends beyond self-improving RAG, GenAI-native agents, and LAMs toward fully collaborative, trustworthy, and ethically aligned Artificial General Intelligence (AGI) systems. The next phase of AI evolution will focus on improving AI-human synergy, ensuring AI systems function as reliable partners rather than opaque decision makers.

By investing in AI safety, promoting responsible AI policies, and fostering global collaboration in AI ethics and governance, we can ensure that the next decade of AI innovation is driven by societal benefit, scientific progress, and technological inclusivity.

The responsibility lies with researchers, policy makers, enterprises, and society as a whole to shape AI's trajectory in a way that maximizes its potential while minimizing its risks. The choices made today will determine

whether AI becomes a transformative force for good or an unchecked system that amplifies biases and vulnerabilities. As AI moves from content generation to real-world execution, the challenge is not just building more powerful AI but ensuring that AI serves humanity in a fair, transparent, and responsible manner.

The year 2025 marks a defining moment in AI's transition from static models to autonomous, action-oriented systems. The convergence of self-improving RAG, GenAI-native agents, and LAMs represents a shift from AI as a passive knowledge generator to AI as an adaptive, decision-making force capable of executing real-world tasks. These advancements will revolutionize industries, scientific research, and societal structures, unlocking unprecedented efficiency, automation, and intelligence. The current state of these technologies varies, with self-improving RAG already seeing substantial adoption, while GenAI-native agents and LAMs remain promising but nascent fields requiring significant breakthroughs, as highlighted in [Table 1](#). The end-to-end workflows and interfaces of these three themes are illustrated in [Figure 2](#).

However, technological, ethical, and security challenges remain, requiring careful oversight to ensure AI development aligns with human values, fairness, and accountability. Issues such as long-term memory, explainability, regulatory compliance, and cybersecurity vulnerabilities must be addressed before AI can be fully trusted in critical domains like healthcare, finance, and governance. The risks associated with autonomous AI

decision making, misinformation, and adversarial threats highlight the urgent need for robust AI governance frameworks.

The responsibility for shaping AI's trajectory does not rest solely on researchers and developers; it extends to governments, enterprises, and society as a whole. Policy makers must establish transparent, enforceable regulations; enterprises must adopt ethical

AI deployment strategies; and the public must be educated on AI's impact and limitations. Ensuring that AI remains a beneficial tool rather than an unchecked force will require collaborative efforts, ethical foresight, and continuous evaluation of AI's evolving role in society.

As AI moves from content creation to real-world execution, the challenge ahead is not just about building more capable AI systems but about

ensuring that AI serves humanity's best interests. The decisions made today, about AI's role in decision making, regulation, and accountability, will determine the future landscape of AI-powered automation.

If implemented responsibly, self-improving RAG, GenAI-native agents, and LAMs will not only enhance productivity and scientific progress but will also redefine how humans and machines

TABLE 1. A summary of key advances in GenAI.					
Technique	Description	Features	Current State	Challenges	Major Takeaways
Self-improving RAG	Real-time adaptive knowledge retrieval integrated with feedback loops	Continuous learning, adaptive indexing, fact verification	Rapid adoption in healthcare, finance, and legal domains	Scalability, misinformation, security vulnerabilities	Enhanced accuracy, real-time knowledge updating
GenAI-native agents	Autonomous agents performing complex reasoning and tasks	Multistep reasoning, task automation, API interactions	Emerging adoption in software development and cybersecurity	Robustness, long-term memory, adversarial attacks	Autonomous problem solving and real-world execution
LAMs	AI capable of autonomously executing tasks in physical and digital environments	Multimodal execution, environmental adaptability, causal planning	Early-stage research and experimentation, limited deployment	Safety validation, uncertainty management, ethical accountability	Action-oriented AI bridging generative content and real-world operations




FIGURE 2. End-to-end workflows and interfaces for emerging GenAI innovations. (a) Self-improving RAG. (b) GenAI-native agents. (c) LAMs

collaborate in an AI-augmented world. The goal for the next decade is clear: to develop AGI systems that are transparent, explainable, secure, and aligned with human values, ensuring that the next generation of AI does not just execute tasks but does so with reliability, fairness, and ethical integrity. 

REFERENCES

1. F. J. G. Peñalvo and A. V. Ingelmo, "What do we mean by GenAI? A systematic mapping of the evolution, trends, and techniques involved in Generative AI," *Int. J. Interactive Multimedia Artif. Intell.*, vol. 8, no. 4, pp. 7–16, 2023.
2. J. Sauvola, S. Tarkoma, M. Klemettinen, J. Riekkki, and D. Doermann, "Future of software development with generative AI," *Automated Softw. Eng.*, vol. 31, no. 1, 2024, Art. no. 26, doi: [10.1007/s10515-024-00426-z](https://doi.org/10.1007/s10515-024-00426-z).
3. M. A. K. Raiaan et al., "A review on large language models: Architectures, applications, taxonomies, open issues and challenges," *IEEE Access*, vol. 12, pp. 26,839–26,874, 2024, doi: [10.1109/ACCESS.2024.3365742](https://doi.org/10.1109/ACCESS.2024.3365742).
4. M. U. Hadi et al., "A survey on large language models: Applications, challenges, limitations, and practical usage," *TechRxiv*, 2023, doi: [10.36227/techrxiv.23589741.v1](https://doi.org/10.36227/techrxiv.23589741.v1).
5. Y. Cao et al., "A comprehensive survey of AI-generated content (AIGC): A history of generative AI from GAN to ChatGPT," 2023, *arXiv:2303.04226*.
6. K. Bellman et al., "Self-improving system integration: Mastering continuous change," *Future Gener. Comput. Syst.*, vol. 117, pp. 29–46, Apr. 2021, doi: [10.1016/j.future.2020.11.019](https://doi.org/10.1016/j.future.2020.11.019).
7. Y. Gao et al., "Retrieval-augmented generation for large language models: A survey," 2023, *arXiv:2312.10997*.
8. S. Siriwardhana, R. Weerasekera, E. Wen, T. Kaluarachchi, R. Rana, and S. Nanayakkara, "Improving the domain adaptation of retrieval augmented generation (RAG) models for open domain question answering," *Trans. Assoc. Comput. Linguistics*, vol. 11, pp. 1–17, Jan. 2023, doi: [10.1162/tacl_a_00530](https://doi.org/10.1162/tacl_a_00530).
9. M. H. Jarrahi, D. Askay, A. Eshraghi, and P. Smith, "Artificial intelligence and knowledge management: A partnership between human and AI," *Bus. Horizons*, vol. 66, no. 1, pp. 87–99, 2023, doi: [10.1016/j.bushor.2022.03.002](https://doi.org/10.1016/j.bushor.2022.03.002).
10. J. Wu, Z. Huang, Z. Hu, and C. Lv, "Toward human-in-the-loop AI: Enhancing deep reinforcement learning via real-time human guidance for autonomous driving," *Engineering*, vol. 21, pp. 75–91, Feb. 2023, doi: [10.1016/j.eng.2022.05.017](https://doi.org/10.1016/j.eng.2022.05.017).
11. J. T. Licardo, M. Domjan, and T. Orehovački, "Intelligent robotics—A systematic review of emerging technologies and trends," *Electronics*, vol. 13, no. 3, 2024, Art. no. 542, doi: [10.3390/electronics13030542](https://doi.org/10.3390/electronics13030542).
12. M. Iftikhar, M. Saqib, M. Zareen, and H. Mumtaz, "Artificial intelligence: Revolutionizing robotic surgery: Review," *Ann. Med. Surgery*, vol. 86, no. 9, pp. 5401–5409, 2024, doi: [10.1097/MS9.0000000000002426](https://doi.org/10.1097/MS9.0000000000002426).
13. C. Zhang et al., "When AI meets finance (StockAgent): Large language model-based stock trading in simulated real-world environments," 2024, *arXiv:2407.18957*.
14. H. N. Fakhouri, B. Alhadidi, K. Omar, S. N. Makhadmeh, F. Hamad, and N. Z. Halalsheh, "AI-driven solutions for social engineering attacks: Detection, prevention, and response," in *Proc. 2nd Int. Conf. Cyber Resilience (ICCR)*, Piscataway, NJ, USA: IEEE Press, 2024, pp. 1–8, doi: [10.1109/ICCR61006.2024.10533010](https://doi.org/10.1109/ICCR61006.2024.10533010).
15. B. M. Mohsen, "AI-driven optimization of urban logistics in smart cities: Integrating autonomous vehicles and IoT for efficient delivery systems," *Sustainability*, vol. 16, no. 24, 2024, Art. no. 11265, doi: [10.3390/su162411265](https://doi.org/10.3390/su162411265).

ABOUT THE AUTHOR

RICHARD SHAN is a consultant and a thought leader in the field of emerging technologies, based in Charlotte, NC 28277 USA. His research interests include enabling organizations to adopt and leverage cutting-edge solutions such as GenAI, LLMOps, agentic systems, RAG, AGI, and more. Contact him at m@richardshan.com.